

# Neural\_Network

December 15, 2019

```
[1]: import matplotlib.pyplot as plt
import numpy as np
from scipy.io import loadmat
x = loadmat('mnist_all.mat')
```

```
[2]: print(x.keys())
```

```
dict_keys(['__header__', '__version__', '__globals__', 'train0', 'test0',
'train1', 'test1', 'train2', 'test2', 'train3', 'test3', 'train4', 'test4',
'train5', 'test5', 'train6', 'test6', 'train7', 'test7', 'train8', 'test8',
'train9', 'test9'])
```

```
[3]: print(len(x['train0']))
print(len(x['train1']))
print(len(x['train0'])+len(x['train1']))
print(len(x['test0']))
print(len(x['test1']))
print(len(x['test0'])+len(x['test1']))
```

```
5923
6742
12665
980
1135
2115
```

```
[4]: Y=np.zeros((1,12665))
```

```
[5]: for i in range(5923,12665):
    Y[0,i]=1
```

```
[6]: X=np.zeros((784,12665))
```

```
[7]: for j in range(12665):
    for i in range(784):
        if j < 5923:
            X[i,j]=x['train0'][j][i]
        else:
```

```
X[i,j]=x['train1'][j-5923][i]
```

```
[8]: X/=255
```

```
[9]: Ytest=np.zeros((1,2115))
for i in range(980,2115):
    Ytest[0,i]=1
Xtest=np.zeros((784,2115))
for j in range(2115):
    for i in range(784):
        if j < 980:
            Xtest[i,j]=x['test0'][j][i]
        else:
            Xtest[i,j]=x['test1'][j-980][i]
Xtest/=255
```

```
[10]: def initialize_parameters_deep(layer_dims):

    parameters = {}
    L = len(layer_dims)      # number of layers in the network

    for l in range(1, L):
        parameters['W' + str(l)] = np.random.randn(layer_dims[l-1],
→layer_dims[l]) * 0.01
        parameters['b' + str(l)] = np.zeros((layer_dims[l], 1))

        assert(parameters['W' + str(l)].shape == (layer_dims[l-1],
→layer_dims[l]))
        assert(parameters['b' + str(l)].shape == (layer_dims[l], 1))

    return parameters
```

```
[11]: def linear_activation_forward(A_prev, W, b, activation):

    if activation == "sigmoid":
        Z, linear_cache = np.dot(W.T,A_prev)+b,(A_prev, W, b)
        A, activation_cache = 1/(1+np.exp(-Z)),Z

    assert (A.shape == (W.T.shape[0], A_prev.shape[1]))
    cache = (linear_cache, activation_cache)

    return A, cache
```

```
[12]: def L_model_forward(X, parameters):

    caches = []
    A = X
    L = len(parameters) // 2      # number of layers in the neural network
```

```

    for l in range(1, L):
        A_prev = A
        A, cache = linear_activation_forward(A_prev, parameters['W' + str(l)],
        ↪parameters['b' + str(l)], "sigmoid")
        caches.append(cache)

    AL, cache = linear_activation_forward(A, parameters['W' + str(L)],
    ↪parameters['b' + str(L)], "sigmoid")
    caches.append(cache)

    assert(AL.shape == (1,X.shape[1]))

    return AL, caches

```

```

[13]: def compute_cost(AL, Y):

    m = Y.shape[1]
    logprobs1 = np.multiply(np.log(AL),Y)
    cost1 = -np.sum(logprobs1)
    logprobs2 = np.multiply(np.log(1-AL),1-Y)
    cost2 = -np.sum(logprobs2)
    cost = (1/m)*(cost1+cost2)

    cost = np.squeeze(cost)
    assert(cost.shape == ())

    return cost

```

```

[14]: def linear_backward(dZ, cache):

    A_prev, W, b = cache
    m = A_prev.shape[1]

    dW = (1/m)*np.dot(A_prev,dZ.T)
    db = (1/m)*np.sum(dZ,axis=1,keepdims=True)
    dA_prev = np.dot(W,dZ)

    assert (dA_prev.shape == A_prev.shape)
    assert (dW.shape == W.shape)
    assert (db.shape == b.shape)

    return dA_prev, dW, db

```

```

[15]: def linear_activation_backward(dA, cache, activation):
    linear_cache, activation_cache = cache

    if activation == "sigmoid":
        A=1/(1+np.exp(-1*activation_cache))

```

```

    dZ = dA*A*(1-A)
    dA_prev, dW, db = linear_backward(dZ, linear_cache)

    return dA_prev, dW, db

```

```

[16]: def L_model_backward(AL, Y, caches):
    grads = {}
    L = len(caches) # the number of layers
    m = AL.shape[1]
    Y = Y.reshape(AL.shape)

    # Initializing the backpropagation
    dAL = - (np.divide(Y, AL) - np.divide(1 - Y, 1 - AL))

    current_cache = caches[L-1]
    grads["dA" + str(L-1)], grads["dW" + str(L)], grads["db" + str(L)] = \
    →linear_activation_backward(dAL, current_cache, activation = "sigmoid")

    for l in reversed(range(L-1)):
        current_cache = caches[l]
        dA_prev_temp, dW_temp, db_temp = linear_activation_backward(grads["dA" + \
    →str(l+1)], current_cache, activation = "sigmoid")
        grads["dA" + str(l)] = dA_prev_temp
        grads["dW" + str(l + 1)] = dW_temp
        grads["db" + str(l + 1)] = db_temp

    return grads

```

```

[17]: def update_parameters(parameters, grads, learning_rate):

    L = len(parameters) // 2 # number of layers in the neural network

    for l in range(L):
        parameters["W" + str(l+1)] = parameters["W" + \
    →str(l+1)] - learning_rate*grads["dW" + str(l+1)]
        parameters["b" + str(l+1)] = parameters["b" + \
    →str(l+1)] - learning_rate*grads["db" + str(l+1)]
    return parameters

```

```

[18]: def L_layer_model(d, X1, Y1, learning_rate = 0.1, num_iterations = 3000, \
    →print_cost=False, plot_graph=False):
    """
    Implements a L-layer neural network: [LINEAR->RELU]*(L-1)->LINEAR->SIGMOID.

    Arguments:
    X -- data, numpy array of shape (num_px * num_px * 3, number of examples)
    Y -- true "label" vector (containing 0 if cat, 1 if non-cat), of shape (1, \
    →number of examples)

```

```

    layers_dims -- list containing the input size and each layer size, of
    →length (number of layers + 1).
    learning_rate -- learning rate of the gradient descent update rule
    num_iterations -- number of iterations of the optimization loop
    print_cost -- if True, it prints the cost every 100 steps

    Returns:
    parameters -- parameters learnt by the model. They can then be used to
    →predict.
    """

    np.random.seed(1)

    costs = []

    layers_dims = [784]
    for i in range(d):
        layers_dims.append(50) #50 nodes in our
    →hidden layer, H=50
    layers_dims.append(1)
    parameters = initialize_parameters_deep(layers_dims)

    # Loop (gradient descent)
    for i in range(0, num_iterations):
        indices = np.random.randint(0,12665,size=100) #We train our model
    →with 100 samples, B=100
        X=X1[:,indices]
        Y=Y1[:,indices]

        AL, caches = L_model_forward(X, parameters)

        cost = compute_cost(AL, Y)

        grads = L_model_backward(AL, Y, caches)

        parameters = update_parameters(parameters, grads, learning_rate)

        # Print the cost every 100 training example
        if print_cost and i % 100 == 0:
            print ("Cost after iteration %i: %f" %(i, cost))
        if print_cost and i % 100 == 0:
            costs.append(cost)

    if plot_graph:
        # plot the cost
        plt.plot(np.squeeze(costs))
        plt.ylabel('cost')


```

```
plt.xlabel('iterations (per hundreds)')
plt.title("Learning rate =" + str(learning_rate))
plt.show()
```

```
return parameters, grads
```

```
[19]: parameters, grads = L_layer_model(2, X, Y, num_iterations = 2500, print_cost = ↵
      ↪True, plot_graph = True)
```

```
Cost after iteration 0: 0.693028
Cost after iteration 100: 0.682784
Cost after iteration 200: 0.673524
Cost after iteration 300: 0.692479
Cost after iteration 400: 0.690433
Cost after iteration 500: 0.687872
Cost after iteration 600: 0.692791
Cost after iteration 700: 0.695646
Cost after iteration 800: 0.696275
Cost after iteration 900: 0.688762
Cost after iteration 1000: 0.693516
Cost after iteration 1100: 0.683833
Cost after iteration 1200: 0.676582
Cost after iteration 1300: 0.579556
Cost after iteration 1400: 0.156000
Cost after iteration 1500: 0.039988
Cost after iteration 1600: 0.030898
Cost after iteration 1700: 0.013863
Cost after iteration 1800: 0.015090
Cost after iteration 1900: 0.007833
Cost after iteration 2000: 0.011290
Cost after iteration 2100: 0.006732
Cost after iteration 2200: 0.003394
Cost after iteration 2300: 0.007008
Cost after iteration 2400: 0.029303
```



```
[20]: def predict(parameters, X):

    # Computes probabilities using forward propagation, and classifies to 0/1,
    →using 0.5 as the threshold.
    A2, cache = L_model_forward(X, parameters)
    predictions = np.zeros((1,X.shape[1]))
    for i in range(A2.shape[1]):
        if A2[0,i] > 0.5:
            predictions[0,i] = 1
        else:
            predictions[0,i] = 0

    return predictions
```

```
[21]: predictions = predict(parameters, X)
print ('Accuracy: %d' % float((np.dot(Y,predictions.T) + np.
    →dot(1-Y,1-predictions.T))/float(Y.size)*100) + '%')
```


Accuracy: 99%

```
[22]: predictions = predict(parameters, Xtest)
print ('Accuracy: %d' % float((np.dot(Ytest,predictions.T) + np.
    →dot(1-Ytest,1-predictions.T))/float(Ytest.size)*100) + '%')
```

Accuracy: 99%

```
[23]: norm = []  
      for d in range(1,11):  
          parameters, grads = L_layer_model(d, X, Y, num_iterations = 2500)  
          norm.append(np.log(np.linalg.norm(grads['dW1'])))
```

```
[24]: x=[1,2,3,4,5,6,7,8,9,10]  
      plt.plot(x,norm)  
      plt.ylabel('Frobenius norm')  
      plt.xlabel('d')  
      plt.title("Graph of log(Frobenius norm) against d")  
      plt.show()
```



```
[25]: def initialize_parameters_deep1(layer_dims):  
  
      parameters = {}  
      L = len(layer_dims)      # number of layers in the network  
  
      for l in range(1, L):  
          parameters['W' + str(l)] = np.random.randn(layer_dims[l-1],  
→layer_dims[l])  
          parameters['b' + str(l)] = np.zeros((layer_dims[l], 1))
```



```

        assert(parameters['W' + str(l)].shape == (layer_dims[l-1],
→layer_dims[l]))
        assert(parameters['b' + str(l)].shape == (layer_dims[l], 1))

    return parameters

```

```

[26]: def L_layer_model_iteration_1(d, X1, Y1, learning_rate = 0.1, num_iterations =
→3000, print_cost=False, plot_graph=False):
    """
    Implements a L-layer neural network: [LINEAR->RELU]*(L-1)->LINEAR->SIGMOID.

    Arguments:
    X -- data, numpy array of shape (num_px * num_px * 3, number of examples)
    Y -- true "label" vector (containing 0 if cat, 1 if non-cat), of shape (1,
→number of examples)
    layers_dims -- list containing the input size and each layer size, of
→length (number of layers + 1).
    learning_rate -- learning rate of the gradient descent update rule
    num_iterations -- number of iterations of the optimization loop
    print_cost -- if True, it prints the cost every 100 steps

    Returns:
    parameters -- parameters learnt by the model. They can then be used to
→predict.
    """

    np.random.seed(1)

    costs = []

    layers_dims = [784]
    for i in range(d):
        layers_dims.append(50)                                #50 nodes in our
→hidden layer, H=50
    layers_dims.append(1)
    parameters = initialize_parameters_deep1(layers_dims)

    indices = np.random.randint(0,12665,size=100)
    X=X1[:,indices]
    Y=Y1[:,indices]

    AL, caches = L_model_forward(X, parameters)

    cost = compute_cost(AL, Y)

    grads = L_model_backward(AL, Y, caches)

```

```

parameters = update_parameters(parameters, grads, learning_rate)

return parameters, grads

```

```

[27]: norm = []
      for d in range(1,11):
          parameters, grads = L_layer_model_iteration_1(d, X, Y, num_iterations = 2500)
          norm.append(np.log(np.linalg.norm(grads['dW1'])))

```

```

[28]: x=[1,2,3,4,5,6,7,8,9,10]
      plt.plot(x,norm)
      plt.ylabel('Frobenius norm')
      plt.xlabel('d')
      plt.title("Graph of log(Frobenius norm) against d")
      plt.show()

```

